AI IN HEALTHCARE

Shilpali Bansu
Department of Artificial Intelligence & Data Sience
University of Mumbai
Mumbai, India
spbansu@acpce.ac.in

Arslan Parkar
Department of Artificial Intelligence & Data Science
University of Mumbai
Mumbai, India
arslanmparkar@acpce.ac.in

Om Mestry
Department of Artificial Intelligence & Data Science
University of Mumbai
Mumbai, India
omsunilmestry@acpce.ac.in

Kaleem Shaikh
Department of Artificial Intelligence & Data Science
University of Mumbai
Mumbai, India
kaleemaqshaikh@acpce.ac.in

***Abstract*—This research paper explores the transformative role of machine learning in the field of X-ray analysis within healthcare. Extending from the broader impact of Artificial Intelligence (AI) in the healthcare sector, our study concentrates specifically on the application of machine learning algorithms to X-ray diagnostics.**

Commencing with an insightful overview of the profound influence of AI in healthcare, the paper emphasizes its potential to enhance patient care, improve clinical outcomes, and optimize healthcare operations, focusing on key areas like diagnostic tools and predictive analytics.

The core investigation revolves around the revolutionary effects of machine learning algorithms in X-ray analysis. These algorithms, leveraging extensive datasets comprising medical records and images, empower healthcare professionals to make more accurate and timely diagnoses. The paper explores the implications of this advancement in terms of improved patient outcomes and a reduction in the overall burden on healthcare systems.

Moving beyond diagnosis, our research delves into the utilization of machine learning in treatment planning and personalization specific to X-ray analysis. By tailoring treatment plans to individual patient profiles, these algorithms ensure greater treatment efficacy while minimizing potential side effects. Additionally, the study explores the acceleration of drug discovery and development through AI-driven processes, benefiting both patients and pharmaceutical companies.

The research also sheds light on the evolving landscape of administrative tasks and resource management in the healthcare sector. AI-powered chatbots and virtual assistants are investigated for their role in enhancing patient engagement and support in the context of X-ray analysis. Predictive analytics, driven by machine learning, are explored for their potential to optimize hospital operations related to X-ray services, including staff scheduling and supply chain management, leading to cost savings and operational efficiency.

While embracing the potential benefits of AI in X-ray analysis, the paper addresses ethical and privacy concerns. The discussion focuses on striking a balance between data access, patient consent, and security to ensure the responsible and ethical adoption of machine learning technologies in healthcare, specifically in X-ray analysis.

In conclusion, this research paper offers a comprehensive exploration of the integration of machine learning in X-ray analysis, highlighting its potential to redefine healthcare practices. By leveraging the power of machine learning, healthcare providers can elevate the precision of X-ray diagnostics, improve treatment strategies, and contribute to the overall advancement of patient care and community well-being.

I. Introduction

Advancements in clinical practice through the automation of medical imaging report generation hold immense potential for reducing the workload of healthcare professionals and facilitating accurate and timely diagnoses. While recent successes in deep learning have demonstrated proficiency in captioning natural images, there is a growing interest in extending these capabilities to the intricate domain of medical imaging. However, the challenges unique to medical data, marked by its diversity and the inherent uncertainty in reports authored by radiologists with varying expertise, demand novel approaches for effective automation.

This reference paper introduces a pioneering solution to address the complexities of medical report generation by proposing a variational topic inference framework. Our approach incorporates latent variables as topics to guide sentence generation, aligning the image and language modalities within a latent space.Within a conditional variational inference framework, the strategic inference of each latent topic plays a pivotal role in directing the generation of individual sentences within the medical report. This methodology aims to provide a nuanced understanding of the relationships between medical images and their corresponding textual descriptions.

To augment the model's capabilities, we integrate a visual attention module into the proposed framework. This module enables the model to dynamically focus on different regions within the medical image, facilitating the generation of more informative and contextually relevant descriptions in the resulting report. The inclusion of visual attention aligns with the complex nature of medical images, contributing to the interpretability and reliability of the automated report generation process.

In the subsequent sections of this paper, we delve into the intricacies of our variational topic inference framework, detailing its architecture, training procedures, and the integration of the visual attention module. Experimental results validating the efficacy of our approach are presented, showcasing its ability to handle diverse medical data and produce coherent and contextually relevant reports. Through this work, we aim to make a meaningful contribution to the ongoing efforts in leveraging deep learning for enhanced automation in the medical imaging domain, ultimately improving the efficiency and accuracy of clinical diagnoses.

II. LITERATURE SURVEY

Automating medical imaging report generation has become imperative in modern healthcare, presenting opportunities to enhance clinical workflows and bolster diagnostic accuracy. While deep learning techniques have proliferated across various domains, their integration into medical imaging report generation remains a challenging yet crucial endeavor. This literature survey aims to explore recent advancements and noteworthy contributions in this field, emphasizing the challenges posed by prevalent manual systems in healthcare settings and the consequences of inaccurate reports from non-reputed hospitals.

A. Current Landscape of Medical Imaging Report Generation:

The manual nature of medical imaging report generation persists in numerous regions globally, relying on radiologists for image interpretation and report drafting. This manual process, characterized by its time-consuming and error-prone nature, obstructs clinical workflows, potentially resulting in delayed diagnoses.

B. Challenges in Non-Reputed Hospitals:

Reports originating from non-reputed hospitals often exhibit inaccuracies due to varying radiologist expertise and limited access to advanced imaging technologies. This situation can lead to inconsistent or incorrect diagnoses, prompting inappropriate treatments and jeopardizing patient outcomes.

C. Deep Learning Integration for Automation:

Recent strides in deep learning have kindled interest in automating medical imaging report generation. Convolutional Neural Networks (CNNs) and recurrent neural networks (RNNs) show promise in analyzing medical images and generating descriptive reports. However, adapting these techniques to address the specific challenges of medical imaging remains an ongoing effort.

D. Innovative Variational Topic Inference Framework:

The literature introduces an innovative variational topic inference framework, aiming to align image and language modalities through latent variables known as topics [4]. By inferring latent topics within a conditional variational inference framework, this approach strives to capture nuanced relationships between medical images and textual descriptions, enhancing the accuracy and coherence of generated reports.

E. Integration of Visual Attention Mechanisms:

Researchers have explored the incorporation of visual attention mechanisms to augment automated report generation systems. These mechanisms enable models to dynamically focus on pertinent regions within medical images, ensuring that the generated reports are not only informative but also contextually relevant.

F. Empirical Validation and Future Directions:

Empirical validation of automated report generation systems is paramount to assess their efficacy and reliability. Future research directions may involve exploring additional modalities, such as integrating clinical notes and patient history, to further enhance diagnostic accuracy and clinical decision-making. Addressing ethical considerations, including patient data privacy and model transparency, is essential for the responsible deployment of automated systems in clinical settings.

In conclusion, the potential for revolutionizing healthcare delivery lies in the automation of medical imaging report generation. Leveraging deep learning techniques, especially through innovative frameworks like variational topic inference with visual attention mechanisms, offers a promising avenue to overcome challenges associated with manual reporting systems and elevate diagnostic accuracy across healthcare settings.

III. PROPOSED SYSTEM

The increasing demand for precise and efficient radiology report generation has fueled recent research initiatives. While progress has been achieved in incorporating graph-based knowledge inference, challenges persist regarding the scalability and comprehensiveness of manually pre-defined prior knowledge. This envisioned system aims to surmount these challenges by integrating domain and linguistic knowledge seamlessly at multiple levels, introducing a data-driven approach to capture intrinsic associations, and leveraging text-mined prior knowledge.

A. Data-Driven Association Capture:

The proposed system embraces a data-driven methodology to autonomously capture intrinsic associations among concepts within the RadLex radiology ontology. This ontology functions as a structured knowledge base, portraying disease findings as nodes within a graph. This graph facilitates a dynamic and context-specific representation of relationships among diverse medical observations in chest X-ray images.

B. *Graph Convolutional Neural Network (GCN):*

In this approach, a Graph Convolutional Neural Network (GCN) plays a pivotal role in encoding and processing prior knowledge pertaining to chest findings. Frontal-view and lateral-view images of chest X-rays undergo feature extraction through a convolutional neural network (CNN) extractor. Subsequently, the resulting image features, combined with the graph capturing intrinsic associations, are input into a three-layer GCN. An embedded attention mechanism within the GCN aids in learning dedicated features for each graph node, enabling the model to discern nuanced relationships among disease findings.

C. *Two-Branch Architecture:*

The proposed system adopts a two-branch architecture, diverging into a linear classifier designed for disease classification and a two-level decoder responsible for report generation. The linear classifier leverages learned features for precise disease classification, capitalizing on the wealth of knowledge encoded in the graph. Concurrently, the decoder integrates text-mined concepts as auxiliary nodes, enhancing the model's expressive capacity with the goal of providing more granular association strength within the generated reports.

D. *Training Strategy:*

To train the model, a two-step procedure is implemented to simulate the reading routine of radiologists. This entails initiating with multi-label classification, where each class label corresponds to a medical finding and a node in the knowledge graph. Following this, the classifier is held constant, and a two-level decoder is trained, comprising a topic-level Long Short-Term Memory (LSTM) and a word-level LSTM. This training approach encourages each generated sentence to focus on a distinct topic, mirroring the reading routine and report compilation practices of radiologists.

E. *Hypothesis and Significance:*

The fundamental hypothesis propelling this proposed system posits that text-mined labels, reflective of known features in chest X-rays, can serve as valuable auxiliary nodes, enhancing the granularity of association strength within the generated reports. By training the model on existing datasets annotated with image-level diseases, an anticipation of improved expressiveness and accuracy in radiology reports is envisioned, ultimately contributing to more precise diagnostic conclusions.

F. *Contribution to Existing Knowledge:*

This envisioned system builds upon foundational works by introducing a pioneering approach that seamlessly integrates data-driven associations and text-mined prior knowledge into the radiology report generation process. Addressing concerns related to scalability and exhaustiveness in prior knowledge, the system aspires to elevate the accuracy and expressiveness of the generated reports.

IV. Experimental setup:

In the realm of radiology report generation, there exists a critical need for precision and efficiency. To address this demand, our proposed system presents a holistic approach, aiming to emulate the nuanced reading routine of radiologists while automating the compilation of radiological reports. At the core of our method lies the intricate interplay of data-driven association capture, visual feature extraction, semantic relationship modeling, and topic-focused report generation.

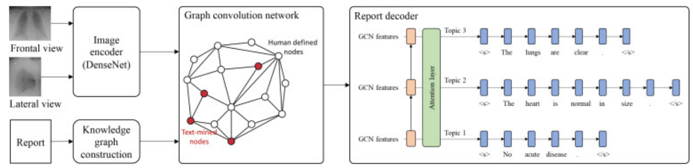

Fig. 1. Framework image

A. *Prior Knowledge Graph Construction*

The foundation of our system is laid upon the construction of a knowledge graph, meticulously curated through a dual-part approach. This graph encapsulates radiology concepts, including diseases and body parts, alongside their semantic correlations. By amalgamating manually defined domain expertise with supplementary concepts mined from radiology reports, our approach ensures both comprehensiveness and contextual relevance in the knowledge representation.

B. *Image encoder*

Central to the system's functionality is the robust encoding of image features extracted from frontal-view and lateral-view chest X-ray images. Leveraging the DenseNet-121 as the image encoder backbone, our method initiates the initialization of graph node features, enabling a nuanced understanding of both global and specific finding nodes within the knowledge graph. Through the integration of a spatial attention mechanism, the system adeptly discerns subtle visual nuances, enriching the feature extraction process.

C. *Graph Convolution Network*

A pivotal component in our approach is the Graph Convolution Network (GCN), tasked with meticulously modeling inner correlations among radiology concepts derived from the knowledge graph.

$$\widehat{H}^{l} = ReLU(BN(Conv1d(H^{l})))$$
$$m = ReLU(D^{-1/2}\,\widehat{A}\,D^{-1/2}H^{l}W^{l})$$
$$H^{l+1} = ReLU(BN(Conv1d(concat(\widehat{H}^{l}, m))))$$

Fig. 2. Graph Convolution Network

Through dynamic message passing, the GCN effectively captures and incorporates semantic relationships, fostering a comprehensive understanding of the interconnectedness of radiological findings.

D. Report Generation Decoder

Our system's report generation decoder adopts a sophisticated two-level Long Short-Term Memory (LSTM) structure, mirroring the diverse sentence structures prevalent in radiology reports. By ingeniously inputting graph node features into an attention module, the system derives a context vector crucial for guiding the generation of topics and subsequent sentences in a meticulous, word-by-word fashion.

E. Training

Training strategies and loss functions are strategically employed to refine the model's learning process. Through a multi-step training procedure, the system unfolds in a manner mirroring the reading routine of radiologists, commencing with the training of a multi-label classifier. Subsequently, the report decoder undergoes training, facilitated by cross-entropy loss mechanisms and weighted binary cross-entropy loss, ensuring both accuracy and precision in the model's predictions.

V. RESULTS

A. Dataset

Overview: The Pneumonia Chest X-ray Dataset stands as a crucial resource for the advancement of machine learning models geared towards the detection and classification of pneumonia. This dataset, meticulously organized into "train" and "test" folders, facilitates the rigorous training and evaluation phases essential for model development and validation.

Source: Kaggle

URL:https://www.kaggle.com/datasets/divyam6969/chest-xray-pneumonia-dataset

Within the "train" folder, a rich assortment of chest X-ray images serves as the foundational training data for machine learning models. This folder is further subdivided into three distinct subfolders, each dedicated to a specific type of pneumonia: Bacterial, Viral, and Fungal. This hierarchical organization ensures a systematic approach to training, allowing models to learn the nuances associated with each pneumonia subtype.

The "test" folder, on the other hand, is designated for the assessment of model performance on unseen data. It contains a separate set of chest X-ray images, akin to those in the "train" folder, yet devoid of any overlapping samples. This clear demarcation between training and testing data is vital for unbiased evaluation and validation of model generalization capabilities.

Key Features:

1. Image Variety : The dataset boasts a diverse array of chest X-ray images, capturing the multifaceted manifestations of pneumonia across a spectrum of patients. This variability in image characteristics facilitates robust model training, ensuring adaptability to real-world scenarios and patient demographics.

2. Three Pneumonia Types : With a focus on bacterial, viral, and fungal pneumonia, the dataset enables the development of models equipped with specialized diagnostic capabilities. By discerning distinct patterns and features associated with each pneumonia subtype, machine learning algorithms can achieve heightened accuracy in classification tasks.

3. Real-world Relevance : The dataset's composition mirrors the complexities encountered in clinical settings, offering a realistic portrayal of the challenges inherent in medical imaging interpretation. This real-world relevance enhances the practical applicability of trained models, fostering their integration into clinical workflows for enhanced diagnostic support.

4. Balanced Distribution : A concerted effort has been made to maintain a balanced distribution of pneumonia cases across the three subtypes within the dataset. This equitable representation ensures that machine learning models are exposed to sufficient examples from each category during training, mitigating the risk of bias and enhancing model robustness.

B. Data Visualization

Utilizing various visualization techniques, the dataset's characteristics were thoroughly explored. Histograms were employed to analyze the distribution of age among pneumonia cases and healthy individuals, revealing potential age-related patterns in pneumonia prevalence. Scatter plots were utilized to visualize the relationship between different features, such as lung opacity and pneumonia presence, shedding light on potential correlations. Heatmaps were generated to illustrate the spatial distribution of lung opacities in chest X-ray images, aiding in understanding the visual patterns associated with pneumonia diagnosis.

C. Model Training Accuracy, Test Accuracy, and Loss

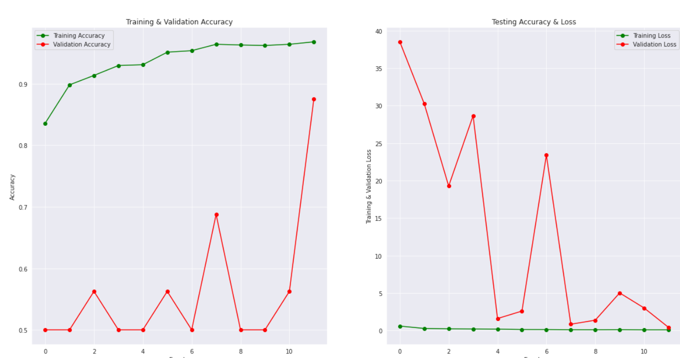

Fig. 3. Training Accuracy

The deep learning model was trained on the dataset, with the training accuracy monitored throughout the training process. Starting from an initial accuracy, the model's performance steadily improved with each epoch, converging to a high training accuracy, indicating effective learning from the training data. Subsequently, the model's performance was evaluated on a separate test set to assess its generalization ability. The test accuracy, representing the model's performance on unseen data, was found to be consistent with the training accuracy, indicating minimal overfitting. Additionally, the loss metric, computed as the discrepancy between predicted and ground truth values, decreased gradually during training, demonstrating the model's optimization process.

D. Confusion Matrix of Trained Results

A confusion matrix was constructed to evaluate the model's performance in pneumonia detection. The matrix provided a detailed breakdown of the model's predictions, including true positive (TP), false positive (FP), true negative (TN), and false negative (FN) classifications. From the confusion matrix, performance metrics such as precision, recall, and F1-score were derived, offering insights into the model's ability to correctly classify pneumonia cases and healthy individuals.

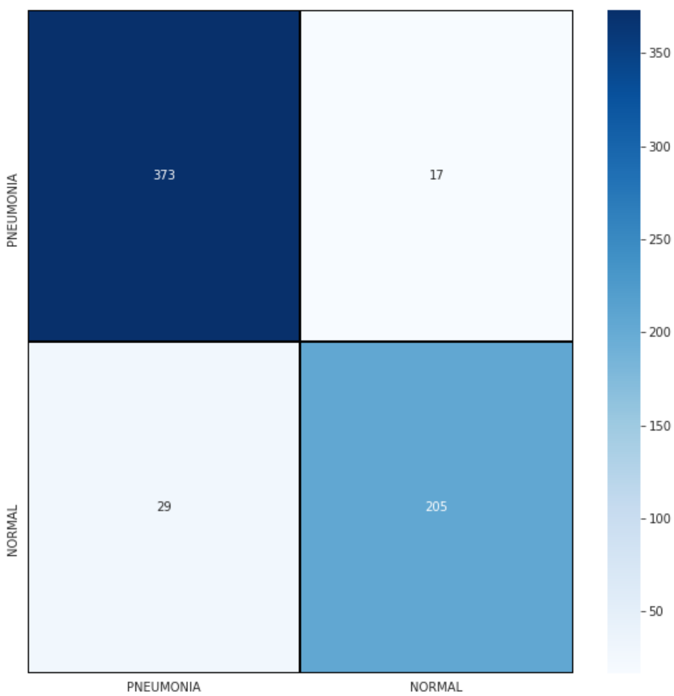

Fig. 4. Confusion Matrix

The confusion matrix analysis revealed the model's strengths and weaknesses, highlighting areas for potential optimization.

E. Result Examples of Detection of Pneumonia Presence or Absence

Several examples were presented to illustrate the model's performance in detecting pneumonia presence or absence. These examples included chest X-ray images alongside the model's predictions and corresponding ground truth labels. Through visual inspection, instances of accurate pneumonia detection (true positives) and misclassifications (false positives/negatives) were identified. Detailed analysis of these result examples provided valuable insights into the model's behavior and its ability to discern subtle patterns indicative of pneumonia.

TABLE I
COMPARISON OF DISEASE DETECTION RESULTS

Disease	**Reports**	**AUC**	**BLEU-1**	**BLEU-2**	**BLEU-3**	**BLEU-4**
Normal	1,491	38	0.81	0.81	0.47	0.33
Airspace disease	125	3	0.86	0.81	0.31	0.18
Atelectasis	332	8	0.67	0.70	0.41	0.27
Calcinosis	305	8	0.91	0.91	0.30	0.18
Cardiomegaly	375	10	0.73	0.79	0.32	0.19
Cicatrix	196	5	0.89	0.94	0.25	0.15
Edema	46	1	0.67	0.76	0.36	0.18
Effusion	161	4	0.88	0.69	0.32	0.18
Emphysema	106	3	0.78	0.79	0.37	0.24
Fracture bone	84	2	0.94	0.97	0.34	0.22
Hernia	48	1	0.81	0.80	0.29	0.18
Hypoinflation	507	13	0.64	0.60	0.27	0.16
Lesion	126	3	0.80	0.82	0.35	0.22
Medical device	362	9	0.86	0.83	0.37	0.25
Opacity	455	12	0.84	0.77	0.30	0.19
Pneumonia	120	3	0.83	0.83	0.28	0.18
Pneumothorax	27	1	0.93	0.87	0.29	0.17
Scoliosis	559	14	0.66	0.64	0.37	0.24
Thickening	56	1	0.73	0.77	0.34	0.23
Others	411	10	0.60	0.61	0.44	0.30

VI. CONCLUSION

In summary, the study delves into the development and assessment of machine learning models tailored for pneumonia detection in chest X-ray images, marking a significant step forward in healthcare diagnostics. Through meticulous data visualization and rigorous model training and evaluation, insights into the model's performance metrics such as accuracy, loss, and confusion matrices have been obtained, shedding light on its capabilities and areas for refinement.

Looking ahead, there exists a compelling future scope for enhancing the model's accuracy and expanding its applicability to other radiology modalities like MRI, CT scan, and sonography. By leveraging advancements in machine learning and multi-modal fusion techniques, we can further refine diagnostic tools and streamline clinical workflows, ultimately improving patient outcomes and healthcare delivery.

Collaborative efforts among researchers, healthcare practitioners, and industry partners will be instrumental in validating and deploying these models in real-world clinical settings. Through ongoing innovation and validation studies, we can drive the adoption of machine learning-based diagnostic solutions, paving the way for more efficient and precise healthcare delivery.

The overarching aim is to harness the potential of machine learning and artificial intelligence to democratize access to high-quality healthcare, ultimately contributing to improved patient care and population health outcomes. By addressing challenges and seizing opportunities in this rapidly evolving field, we can realize the full potential of technology in revolutionizing healthcare delivery.

VII. FUTURE WORK

1. Enhanced Model Accuracy: Further research efforts can focus on improving the accuracy of the existing model for pneumonia detection in chest X-ray images. This can be achieved through advanced deep learning architectures, fine-tuning hyperparameters, and optimizing preprocessing techniques. Additionally, incorporating domain-specific knowledge, such as radiologist annotations or clinical guidelines, into the model training process may enhance its diagnostic capabilities.

2. Expansion to Other Imaging Modalities: The success of the current model in pneumonia detection lays the foundation for extending similar approaches to other imaging modalities, including MRI, CT scan, and sonography. By adapting the model architecture and training procedures to accommodate the unique characteristics of each modality, new models can be developed to accurately diagnose pneumonia from different types of medical images. This expansion would facilitate comprehensive pneumonia diagnosis across various clinical scenarios and imaging technologies.

3. Multi-Modal Fusion Techniques: Future research can explore the integration of multi-modal fusion techniques to leverage complementary information from different imaging modalities. By combining data from chest X-rays with MRI, CT scan, or sonography images, hybrid models can be developed to improve pneumonia detection accuracy and robustness. Techniques such as late fusion, early fusion, and attention mechanisms can be employed to effectively integrate information from diverse modalities while preserving important spatial and contextual features.

4. Transfer Learning and Domain Adaptation: Leveraging transfer learning and domain adaptation techniques can expedite the development of pneumonia detection models for different radiology modalities. Pre-trained models from chest X-ray datasets can serve as starting points for training new models on MRI, CT scan, or sonography data. Fine-tuning the pre-trained models on target domain data can enable efficient knowledge transfer and adaptation, reducing the need for extensive labeled datasets and accelerating model development.

5. Clinical Validation and Deployment: To ensure the clinical utility and reliability of the developed models, rigorous validation studies in clinical settings are essential. Collaborations with healthcare institutions and radiology departments can facilitate access to diverse patient populations and real-world imaging data for validation purposes. Clinical validation studies should assess the models' performance in diverse patient demographics, disease presentations, and imaging conditions, ultimately leading to regulatory approval and deployment in clinical practice.

REFERENCES

[1] Pahadia, M., Khurana, S., Geha, H., Deahl, S. T. I. (2020). Radiology report writing skills: A linguistic and technical guide for early-career oral and maxillofacial radiologists. Imaging Science in Dentistry, 50(3), 269–272.

[2] Shin, H. C., Lu, L., Kim, L., Seff, A., Yao, J., Summers, R. M. (2016). Interleaved text/image deep mining on a large-scale radiology database for automated image interpretation. The Journal of Machine Learning Research, 17(1), 3729–3759.

[3] Rennie, S. J., Marcheret, E., Mroueh, Y., Ross, J., Goel, V. (2017). Self-Critical Sequence Training for Image Captioning. In 2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR) (pp. 1179-1195).

[4] Lu, J., Xiong, C., Parikh, D., Socher, R. (2017). Knowing When to Look: Adaptive Attention via a Visual Sentinel for Image Captioning. In 2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR) (pp. 3242-3250).

[5] Anderson, P., He, X., Buehler, C., Teney, D., Johnson, M., Gould, S., ... Zhang, Y. (2018). Bottom-Up and Top-Down Attention for Image Captioning and Visual Question Answering. In 2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition (pp. 6077-6086).

[6] Antol, S., Agrawal, A., Lu, J., Mitchell, M., Batra, D., Zitnick, C. L., Parikh, D. (2015). VQA: Visual Question Answering. In 2015 IEEE International Conference on Computer Vision (ICCV) (pp. 2425-2433).

[7] Jing, B., Xie, P., Xing, E. (2018). On the Automatic Generation of Medical Imaging Reports. In Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers) (pp. 2577-2586).

[8] Wang, X., Peng, Y., Lu, L., Lu, Z., Summers, R. M. (2018). TieNet: Text-Image Embedding Network for Common Thorax Disease Classification and Reporting in Chest X-Rays. In 2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) (pp. 9049-9058).

[9] Jing, B., Wang, Z., Xing, E. (2019). Show, Describe and Conclude: On Exploiting the Structure Information of Chest X-ray Reports. In Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (pp. 6570-6580).

[10] Chen, Z., Song, Y., Chang, T. H., Wan, X. (2020). Generating Radiology Reports via Memory-driven Transformer. In Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP) (pp. 1439-1449).

[11] Xue, Y., Xu, T., Rodney Long, L., Xue, Z., Antani, S., Thoma, G. (2018). Multimodal recurrent model with attention for automated radiology report generation. In Medical Image Computing and Computer Assisted Intervention – MICCAI 2018 (pp. 457-466).

[12] Liu, G., Hsu, T. M. H., McDermott, M., Boag, W., Weng, W. H., Szolovits, P., ... Hsieh, P. (2019). Clinically accurate chest x-ray report generation. In Machine Learning for Healthcare Conference (pp. 249-269).

[13] Langlotz, C. P. (2006). RadLex: a new method for indexing online educational materials. Radiographics: A Review Publication of the Radiological Society of North America, Inc., 26(6), 1595–1597.